


# GOSSIP-STYLE FAILURE DETECTION



## Protocol

- Nodes periodically gossip their membership list
- On receipt, the local membership list is updated

Current time: 70 at node 2  
(asynchronous clocks)

Fig and animation by: Dongyun Jin and Thuy Ngyuen




# GOSSIP-STYLE FAILURE DETECTION

- If the heartbeat has not increased for more than  $T_{\text{fail}}$  seconds, the member is considered failed
- And after  $T_{\text{cleanup}}$  seconds, it will delete the member from the list
- Why two different timeouts?

# GOSSIP-STYLE FAILURE DETECTION




- What if an entry pointed to a failed node is deleted right after  $T_{\text{fail}}$  (=24) seconds?



# GOSSIP-STYLE FAILURE DETECTION




- What if an entry pointed to a failed node is deleted right after  $T_{\text{fail}}$  ( $=24$ ) seconds?



# GOSSIP-STYLE FAILURE DETECTION




- What if an entry pointed to a failed node is deleted right after  $T_{\text{fail}}$  ( $=24$ ) seconds?



# GOSSIP-STYLE FAILURE DETECTION




- What if an entry pointed to a failed node is deleted right after  $T_{\text{fail}}$  (=24) seconds?



# GOSSIP-STYLE FAILURE DETECTION



- What if an entry pointed to a failed node is deleted right after  $T_{\text{fail}}$  (=24) seconds?



- Fix: remember for another  $T_{\text{fail}}$